

EECS 489 Lab 2: A Peer Node

This assignment is due on **Wednesday, 21 Jan 2015, 10 pm.**

Introduction

The majority of socket programs, including Netimg of Lab1, follows the client-server paradigm, where a server waits on a well-known port for clients' connections. In this lab, we'll explore peer-to-peer programming. A peer is basically both a server and a client. It accepts connections from other peers and also connects to one or more peers.

You're provided with a skeleton code, a single file called `peer.cpp` as part of this lab. You can download the [support code](#) from the Course Folder. The provided `Makefile` builds a program called `peer`. The program takes a single, optional argument on the command line:

```
% peer [-p <hostname>:<port>]
```

The `-p` option tells the `peer` program which peer to connect to initially. If this option is not provided, the `peer` starts as a server listening on a random, ephemeral port.

To boot strap the peer-to-peer (p2p) network, we first start a peer by itself. Everytime a peer runs, it prints to screen/console its fully qualified domain name (FQDN) and the port number it is listening on. When a peer is run with the `hostname:port` of another peer as its command line argument, the new peer tries to join the provided peer in the p2p network by creating a socket and connecting to the peer.

A peer that receives a join request will accept the peer if and only if its peer table is not full. Whether a join request is accepted or not, the peer sends back to the requesting peer the `hostname:port` of a peer in its peer table, if the table is not empty, to help the newly joined peer find more peers to join.

In completing this lab, you may consult the sample code `server.c` and `client.c` from Lab 0 and your code for Lab 1. In terms of code you have to write, the majority of it is very similar to what you did in Lab 1. Try to take a step back and look at the big picture, i.e., how the two pieces of code that were implemented in two different processes now reside in the same process and how this process is serving the role of both a client and a server. Pay particular attention to how this is accomplished by monitoring multiple sockets on a single thread. Another goal of this lab is for you to gain an early experience with protocol design. In this case, we're designing a simple peer-to-peer join protocol, with redirection.

Task 1: Server Side

Your first task is to implement the server side of a peer. You can search for the string "Task 1" in the code to find places where "Task 1" related code must be filled in. You can search for the string "YOUR CODE HERE" in the code to find places where your code must go.

If `peer` is run without any option on the command line, it calls `peer_setup(port)` with the argument `port = 0`. Fill in the function `peer_setup(port)` by first creating a TCP socket. Since we will be re-using the same port number with both a server listening socket and a client connect socket, set the socket option to allow for address reuse. Then bind the socket to the port passed in as argument to `peer_setup()` and listen for connection, with listen queue length set to the macro `PR_QLEN`. If `port = 0`, the OS will assign a

random, ephemeral port to the socket. Finally, return the socket descriptor to the caller. The `peer_setup()` has been commented such that it should be clear where you need to make which socket API call. Depending on the error checking you do, it takes only 5 to 9 lines of code to complete this function.

Back in `main()`, if `peer` was run without any command line option, `peer_setup()` would have obtained a random, ephemeral port number on the returned socket. Find out the port number assigned to the socket and store it in the `self` variable. Next find out the name of the host the peer is currently running on. Store the name in the memory space pointed to by `pname[1]`, which we're using as scratch space. The current host name is used for printing user-friendly status messages to the console. This part of the task takes 3 to 5 lines of code.

Next call `select()` to wait for connection on the listened on socket (1 to 2 lines of code). When a connection is made, `main()` checks if its peer table is full (for this lab, we restrict the peer table size to 2). If the table is not full, `main()` calls `peer_accept()` to accept the connection and then calls `peer_ack()` to send back a welcome (`PM_WELCOME`) message with a list of all the other peers it knows of, if any (well, in this lab, "all" means the only other peer in its peer table). The new peer is then stored in the peer table. On the other hand, if the peer table is full, `main()` calls `peer_accept()` and `peer_ack()` as before, but in the call to `peer_ack()`, it sends back a redirect (`PM_REDIRECT`) message, along with the first peer it knows of, if any, and closes the connection.

The function `peer_accept(sd, pte)` accepts the connection on the socket `sd`. Since we will be sending back acknowledgement message when the peer table is full and we must close the connection, set the socket option so that the socket will linger for `PR_LINGER` amount of time upon closing to give time for the acknowledgement message to arrive at the redirected peer. This part takes 5 to 7 lines of code.

The function `peer_ack(td, type, peer)` marshalls together a message of type `pmsg_t` defined at the top of `peer.cpp`. It fills in the fields of the message: `pm_vers` must be set to `PM_VERS`, `pm_type` set to the type argument passed in to `peer_ack()`, if the `peer` pointer passed in is a `NULL` pointer, set `pm_npeers` to 0, otherwise set it to 1 and copy the address and port of the peer pointed to by `peer` to `pm_peer`. Then `peer_ack()` sends the marshalled message through the provided socket `td`. If there's any error in sending, for example, if the other side of the connection has been closed by the peer, close the connection. This part takes about 13 lines of code.

That's all for Task 1. It should take about 27 to 36 lines of code in total. After completing Task 1, you should test your code before continuing to Task 2. See the Testing section below for some guidelines on testing your code using the reference implementation of `peer`.

Task 2: Client Side

You can search for the string "Task 2" in the code to find places where "Task 2" related code must be filled in.

If `peer` is run with the `-p` option, the user must provide a peer hostname and a port number of the peer to connect to, with the port number separated from the peer hostname by a colon. The provided function `peer_args()` handles parsing of the command line. Upon return from the call to `peer_args()`, the peer's hostname will be stored in `pname[0]` and the port number will be stored, in network byte order, in `pte[0]`. Given the peer's hostname stored in `pname[0]`, determine the peer's IPv4 address and store it in peer table entry 0 (`pte[0]`). Then call the `peer_connect()` function, passing it a pointer to the first peer table entry. When `peer_connect()` connects to the provided peer, it will be assigned a random, ephemeral port number by the OS.

The function `peer_connect(pte)` connects to the provided peer. First create a new TCP socket, store the new socket descriptor in `pte->pte_sd`. Since we will be re-using the same port number with both a server listening socket and a client connect socket, set the socket option to allow for address reuse. Next initialize the socket with the destination peer's IPv4 address and port number with the peer table entry pointed to by the `pte` argument. Finally connect to the destination peer and return to caller. When you connect to the destination peer, the OS will assign a random, ephemeral port to the socket. If there were any error during the connect process, terminate process.

Back in `main()`, find out the assigned ephemeral port number and store it in the `self` variable, along with the IPv4 address of the current host. The function `peer_connect()` should take 5 to 10 lines of code. The code in `main()` prior to and upon return from the call to `peer_connect()` together should take about 4 to 6 lines.

At this point in `main()`, we'll be calling the `peer_setup()` function you wrote earlier in Task 1. However, instead of calling the function with `port = 0`, we'll be calling it with the random, ephemeral port number assigned by the OS when you connected with the user-provided peer. In the call to `select()`, we will be waiting for activities on the listened to socket and all connected socket(s), if any. When a message arrives from a connected peer, we call `peer_rcv(td, msg)`.

The function `peer_rcv(td, msg)` receives a message from the provided socket `td` into the buffer space pointed to by the provided `msg` pointer, and returns the size of the received message, which in this lab should always be `sizeof(pmsg_t)`. If there is any error in receiving the message, close the socket `td` and return to caller the error code returned by the socket receive API. This function should take about 11 lines of code. Back in `main()`, if the received packet contains another peer, print out the third peer's address and port number. If the received packet is a `PM_RDIRECT` packet, inform the user that the join has been declined (redirected). The user can manually try to connect to the third peer returned in the redirect packet.

That's all for Task 2. The total number of lines for Task 2 should be about 20 to 27 lines of code. And the total number of lines for both tasks together is about 47 to 62 lines.

Testing Your Code

We will use the same four hosts CAEN has set up for this course. Again, don't use CAEN's login server (`login.engin.umich.edu`) which will redirect you to one of `caen-vnc*` hosts as these hosts do not allow for connection to random ports. You can also run multiple peers on a single host and form p2p connections between them. When multiple peers are running on the same host, you can use `localhost` in place of the peer's hostname in the command line to peer.

In addition to the skeletal code and `Makefile`, we've also provided an executable binary of peer, called `refpeer`, that runs on CAEN `eecs489` hosts. It is available on `/afs/umich.edu/class/eecs489/w15/FILES/`. As in Lab 1, this is a GNU/Linux executable, not to be downloaded nor run on your Mac OS X, Ubuntu, nor Windows machines. Remember that you can connect to the CAEN `eecs489` hosts only through [UMVPN](#) and MWireless or from CAEN Lab desktops. You should test your code as soon as you completed Task 1. Use `refpeer` to connect to your peer. Similarly, after completing Task 2, connect your peer to `refpeer`. To see the expected behavior of the code, run multiple `refpeers` and have them connect to each other.

Here is an example test scenario, assuming that you have built the program `peer` and it is residing in your working directory/folder for this lab. Create four windows on your local host.

1. On the first window, ssh to `eeecs489p1.engin.umich.edu`, change to your working directory for this lab, run `peer` without any command line argument:

```
p1% ./peer
```

It should print to screen (with a different port number, depicted in bold here):

```
This peer address is caen-eeecs489p01.engin.umich.edu:43945.
```

Note that `eeecs489p1.engin.umich.edu` is an alias/CNAME for `caen-eeecs489p01.engin.umich.edu`. On the four `eeecs489` machines, but not from your laptop, you can also refer to each of them as `p1` to `p4`.

2. On the second window, ssh to `eeecs489p2.umich.edu`, change to your working directory for this lab, run `peer` with the following command line argument (replacing the port number with the one that got printed for you on the first item above):

```
p2% ./peer -p p1:43945
```

It should print to screen (with different port numbers):

```
Connected to peer p1:43945
This peer address is caen-eeecs489p02.engin.umich.edu:56535
Received ack from p1:43945
```

Meanwhile, on the first window, you should see the following additional line printed to screen:

```
Connected from peer p2:56535
```

3. On the third window, ssh to `eeecs489p3.engin.umich.edu`, change to your working directory for this lab, run `peer` with the following command line argument (replacing the port number with the one from the first item above):

```
p3% ./peer -p p1:43945
```

It should print to screen (with different port numbers):

```
Connected to peer p1:43945
This peer address is caen-eeecs489p02.engin.umich.edu:48141
Received ack from p1:43435
  which is peered with: p2:56535
```

Meanwhile, on the first window, you should see the following additional line printed to screen:

```
Connected from peer p3:48141
```

4. On the fourth window, ssh to `eeecs489p4.engin.umich.edu`, change to your working directory for this lab, run `peer` with the following command line argument (replacing the port number with the one from the first item above):

```
p4% ./peer -p p1:43945
```

It should print to screen (with different port numbers):

```
Connected to peer p1:43945
```

```
This peer address is caen-eecs489p04.engin.umich.edu:40231
Received ack from p1:43945
  which is peered with: p2:56535
Join redirected, try to connect to the peer above.
```

Meanwhile, on the first window, you should see the following additional line printed to screen:

```
Peer table full: p4:40231 redirected
```

5. Staying on the fourth window, if the peer hasn't automatically exited, terminate it by entering 'q' (and hit "enter"), and run peer again with the following command line argument (replacing the port number with the one from the **fourth** item above):

```
p4% ./peer -p p2:56535
```

It should print to screen (with different port numbers):

```
Connected to peer p2:56535
This peer address is caen-eecs489p04.engin.umich.edu:50095
Received ack from p2:56535
  which is peered with: p1:43945
```

Meanwhile, on the second window, on p2, you should see the following additional line printed to screen:

```
Connected from peer p4:50095
```

That ends our sample test scenario and you can quit all four peers.

Remember not to use any libraries or compiler options not already used in the Makefile to ensure that we will be able to build your code for grading. If we can't compile your code, you will get 0 point.

Submission Instructions

Test your compilation! Your submission must compile **without** errors.

Your "*Lab2 files*" comprises your `peer.cpp` file only.

To turn in your Lab2:

1. Email your IA/GSI the SHA1's of your *Lab2 file*. Use "EECS489: Lab2 Submission" as your email's "Subject:" line. Once you've sent in your SHA1's, **don't make any more changes to the files**, or your SHA1's will become invalid.
2. Upload your *Lab2 files* by pointing your web browser to [Course folder](#) and navigate to your lab2 folder under your username. Or you can scp the files to your lab1 folder on IFS:
`/afs/umich.edu/class/eecs487/w15/FOLDERS/<username>/lab2/`
This path is accessible from any machine you've logged into using your ITCS (umich.edu) password. Please report any problems to ITCS.
3. **Keep your own backup copy!** Don't make any more changes to the files once you've submitted your final SHA1's.

The timestamp on your SHA1 email will be your time of submission. If this is past the deadline, your submission will be considered late. You are allowed multiple "submissions" without late-policy

implications as long as you respect the deadline. **Try not to email your SHA1 to your IA/GSI until you've finalized your code.** You don't want to annoy them.

Do NOT turn in an archival (.zip or .tgz) file, instead please turn in your solution files individually. **Turn in ONLY the files you have modified.** Do not turn in support code we provided that you haven't modified. **Do not turn in any binary files (object files, executables, or images) with your assignment.**

Do **remove** all `printf()`'s or `cout`'s and `cerr`'s and any other logging statements you've added for debugging purposes. You should debug using a debugger, not with `printf()`'s. If we can't understand the output of your code, you will get zero point. You can keep error reporting messages that you print out prior to terminating your code.